

PolicyEngine Atlas Project Roadmap

Phase 1: Scale (Months 1-3)

Goal: Expand from 237 to 2,500 documents

Month 1: Infrastructure Setup

- **Week 1-2:** Set up core infrastructure on Google Cloud Platform
 - Configure Git repositories with LFS for document storage
 - Deploy initial API endpoints
 - Establish GitHub PR review workflow
- **Week 3-4:** Develop AI crawler framework
 - Implement LLM-based document extraction
 - Create metadata standardization system
 - Build initial validation pipeline

Deliverable: Working crawler system with 500 documents processed

Month 2: PolicyEngine Citation Migration

- **Week 1-2:** Extract and validate 2,500 existing PolicyEngine citations
 - Parse URLs from rules engine
 - Download and archive documents
 - Generate initial metadata
- **Week 3-4:** Launch first bounty program
 - Partner with Urban Institute and GCO for validation
 - \$5,000 allocated for metadata verification
 - Establish quality control metrics

Deliverable: 1,500 documents archived with validated metadata

Month 3: Partner Integration Begins

- **Week 1-2:** Deploy initial API for partners
 - MyFriendBen integration for Colorado/North Carolina
 - Benefit Navigator pilot in LA County
 - Atlanta Fed Policy Rules Database connection
- **Week 3-4:** Implement feedback and iterate
 - Refine metadata schema based on partner needs
 - Optimize API performance
 - Document integration patterns

Deliverable: 2,500 documents accessible via API with 3+ partner integrations

Phase 2: Monitor & Update (Months 4-9)

Goal: Reach 5,000 documents with change tracking

Month 4-5: Automated Monitoring System

- Develop change detection algorithms

- Implement diff generation for policy updates
- Create notification system for stakeholders
- Replace direct URLs with Atlas permalinks

Deliverable: Monitoring system tracking 3,000+ documents

Month 6-7: Clarity Index Development

- Design scoring methodology with human baseline
- Implement LLM consistency testing
- Validate against SNAP QC error data
- Create jurisdiction-level dashboards

Deliverable: Clarity Index scores for 1,000 priority documents

Month 8-9: Semantic Layer & Search

- Build knowledge graph connecting documents to rules
- Implement semantic search using embeddings
- Create categorical eligibility pathway detection
- Launch web interface for caseworker access

Deliverable: Searchable archive of 4,500 documents with semantic relationships

Phase 3: Discover & Complete (Months 10-12)

Goal: Comprehensive coverage with community validation

Month 10: AI Discovery Launch

- Deploy AI agents to find missing documents
- Focus on operational manuals and memos
- Identify gaps in current coverage
- Second bounty program (\$7,500) for verification

Deliverable: 500 new documents discovered and verified

Month 11: Community Contribution Platform

- Launch third bounty program (\$7,500) for gap filling
- Enable direct agency submissions
- Create contributor recognition system
- Expand to all 50 states plus territories

Deliverable: 5,000 total documents with nationwide coverage

Month 12: Production Deployment & Evaluation

- Complete LLM benchmark study
- Publish impact assessment
- Formalize government partnerships
- Transition to sustainable operations

Deliverable: Full production system with documented impact

Year 2: Sustainability & Scale (Months 13-24)

Months 13-18: Advanced Features

- MCP server for AI assistant integration
- Advanced analytics and reporting
- Historical document versioning
- Enhanced metadata and tagging systems

Months 19-24: Ecosystem Development

- Expand to additional program types
- Develop training materials and documentation
- Build community of practice
- Establish revenue streams for sustainability

Key Milestones & Success Metrics

Milestone	Target Date	Success Metric
Infrastructure Live	Month 1	Crawler processing 100 docs/day
Initial Scale	Month 3	2,500 documents archived
Partner Adoption	Month 6	10+ organizations integrated
Nationwide Coverage	Month 11	All 50 states represented
Production Launch	Month 12	5,000 documents, 1000s of API calls/month
Sustainability	Month 24	Revenue covering operational costs

Risk Mitigation

Technical Risks

- **Crawler failures:** Manual fallback processes, multiple extraction methods
- **Storage scaling:** Cloud architecture designed for 100,000+ documents
- **API performance:** Caching, CDN, and rate limiting strategies

Operational Risks

- **Partner delays:** Staggered rollout, flexible integration timeline
- **Document access:** Multiple sourcing strategies, FOIA backup plans
- **Quality control:** Multi-tier review process, automated validation

Financial Risks

- **Cost overruns:** Conservative cloud estimates, reserved capacity
- **Bounty program:** Fixed budgets with clear deliverables
- **Sustainability:** Multiple revenue streams identified